

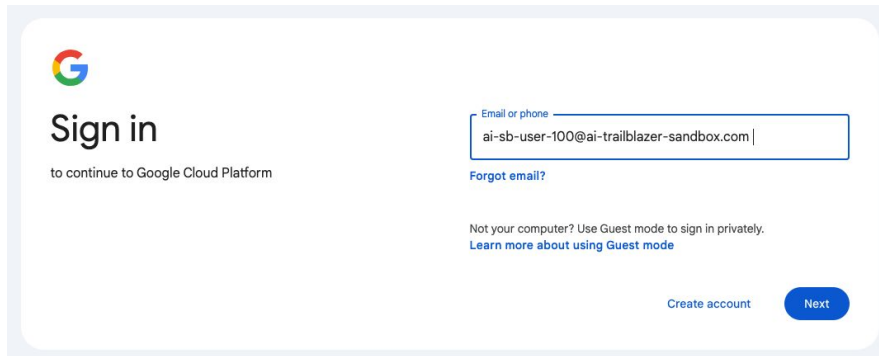
# LAB GUIDE

AI Labs - Prudential with Google

# Labs Access: Accessing Google Cloud Console

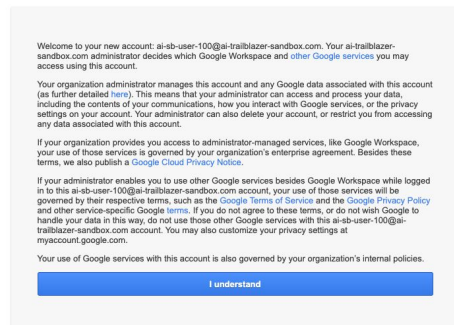
**Step 1:** Navigate to Google Cloud Console.

<https://console.cloud.google.com/>



The screenshot shows the Google Cloud Platform sign-in interface. At the top left is the Google logo. Below it, the text "Sign in" is displayed, followed by "to continue to Google Cloud Platform". To the right, there is a text input field labeled "Email or phone" containing the email address "ai-sb-user-100@ai-trailblazer-sandbox.com". Below the input field are two links: "Forgot email?" and "Not your computer? Use Guest mode to sign in privately. [Learn more about using Guest mode](#)". At the bottom right, there are two buttons: "Create account" and "Next".

**Step 2:** For the first login, you have to click “I Understand”



The screenshot shows the "Welcome to your new account" page. At the top, the Google logo is displayed, followed by the text "Welcome to your new account". Below this, there is a large block of text explaining the account setup and terms of service. At the bottom, there is a blue button labeled "I understand".

Welcome to your new account: ai-sb-user-100@ai-trailblazer-sandbox.com. Your ai-trailblazer-sandbox.com administrator decides which Google Workspace and other Google services you may access using this account.

Your organization administrator manages this account and any Google data associated with this account (as further detailed [here](#)). This means that your administrator can access and process your data, including the contents of your communications, how you interact with Google services, or the privacy settings on your account. Your administrator can also delete your account, or restrict you from accessing any data associated with this account.

If your organization provides you access to administrator-managed services, like Google Workspace, your use of those services is governed by your organization's enterprise agreement. Besides these terms, we also publish a [Google Cloud Privacy Notice](#).

If your administrator enables you to use other Google services besides Google Workspace while logged in to this ai-sb-user-100@ai-trailblazer-sandbox.com account, your use of those services will be governed by their respective terms, such as the [Google Terms of Service](#) and the [Google Privacy Policy](#) and other service-specific Google terms. If you do not agree to these terms, or do not wish Google to handle your data in this way, do not use those other Google services with this ai-sb-user-100@ai-trailblazer-sandbox.com account. You may also customize your privacy settings at [myaccount.google.com](#).

Your use of Google services with this account is also governed by your organization's internal policies.

[I understand](#)

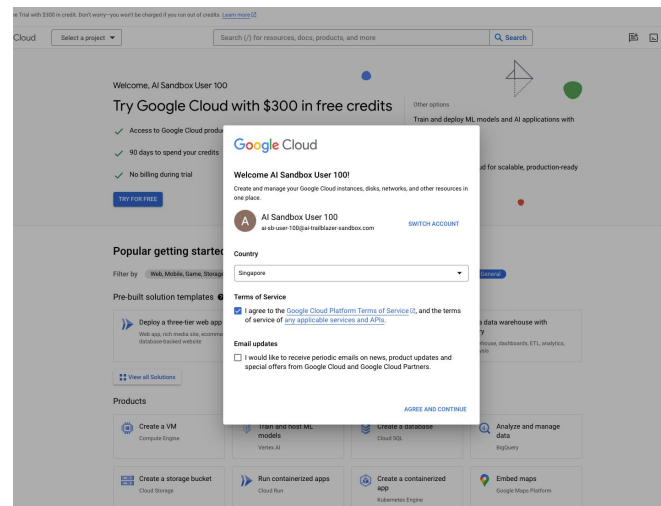
# Labs Access: Accessing Google Cloud Console

Proprietary + Confidential

**Step 3:** You will be prompted to change password. Please enter a new password.



**Step 4:** Check the checkbox for “Terms of Service” and then proceed to click on “Agree and Continue”.



# Labs Access: Accessing Google Cloud Console

Proprietary + Confidential

## Navigate to Vertex AI and then Workbench

The screenshot displays the Google Cloud console interface. In the left-hand navigation menu, the 'Vertex AI' option is highlighted with a red rectangular box. Below it, under the 'NOTEBOOKS' section, the 'Workbench' option is also highlighted with a red rectangular box. The main content area shows the 'Workbench' page, which includes tabs for 'INSTANCES', 'EXECUTIONS', and 'SCHEDULES'. The 'INSTANCES' tab is active, showing a table of Workbench instances. The table has columns for Instance name, Zone, Auto upgrade, Version, Machine Type, GPUs, and Owner. Three instances are listed, each with a green checkmark in the status column and a link to 'OPEN JUPYTERLAB'.

Instance name	Zone	Auto upgrade	Version	Machine Type	GPUs	Owner
<a href="#">ai-sb-user-523-gpu</a>	asia-southeast1-a	—	M124	Graphics Optimized: 1 NVIDIA L4 GPU, 4 vCPUs, 16GB RAM	NVIDIA L4 x 1	391264778563-compute@developer.gserviceaccount.c
<a href="#">ai-sb-user-523-no-gpu</a>	asia-southeast1-a	—	M124	Efficient Instance: 4 vCPUs, 16 GB RAM	None	391264778563-compute@developer.gserviceaccount.c
<a href="#">ai-sb-user-524-gpu</a>	asia-southeast1-a	—	M124	Graphics Optimized: 1 NVIDIA L4 GPU, 4 vCPUs, 16GB RAM	NVIDIA L4 x 1	391264778563-compute@developer.gserviceaccount.c

# Labs Setup: Accessing Workbench Notebook

Proprietary + Confidential

**Step 4:** Once at Workbench, ensure you are on the **INSTANCES** tab

**Step 5:** You should see that the notebooks created for your team.  
Each user is assigned 2 notebooks. 1 with GPU and another without GPU.

GPU instance: ai-gpu-100-XXX. Non-GPU instance: ai-no-gpu-100--XXX

Google Cloud | ai-sandbox-Company17 | Search (/) for resources, docs, products, and more | Search

Vertex AI | Workbench | CREATE NEW | REFRESH | LEARN

TOOLS: Dashboard, Model Garden, Pipelines

NOTEBOOKS: Colab Enterprise, Workbench

VERTEX AI STUDIO: Overview, Multimodal, Language, Vision, Speech

DATA

View: **INSTANCES** | USER-MANAGED NOTEBOOKS | MANAGED NOTEBOOKS

Workbench Instances have JupyterLab 3 pre-installed and are configured with GPU-enabled machine learning frameworks. [Learn more](#)

Filter	Instance name	Zone	Auto upgrade	Version	Machine Type	GPUs	Owner	Created	Labels
<input checked="" type="checkbox"/>	<a href="#">ai-gpu-100-3ffb19a7</a> OPEN JUPYTERLAB	asia-southeast1-c	—	M118	Graphics Optimized: 1 NVIDIA L4 GPU, 4 vCPUs, 16GB RAM	NVIDIA L4 x 1	55282892569-compute@developer.gserviceaccount.com	Mar 11, 2024, 7:49:27 AM	consumer.p...: ai-sandbox...
<input checked="" type="checkbox"/>	<a href="#">ai-sb-user-100-no-gpu</a> OPEN JUPYTERLAB	asia-southeast1-c	—	M118	Efficient Instance: 4 vCPUs, 16 GB RAM	None	55282892569-compute@developer.gserviceaccount.com	Mar 11, 2024, 7:56:19 AM	consumer.p...: ai-sandbox...

# Labs Setup: Accessing Workbench Notebook

**Step 6:** Select the 2 instances assigned and click on **START** in the top menu. If the instances are already started, skip to step 8.

The screenshot shows the Google Cloud Vertex AI Workbench interface. At the top, there's a header with the Google Cloud logo, a dropdown menu for 'ai-sandbox-Company17', a search bar, and buttons for 'DISMISS' and 'START FREE'. Below the header, a blue banner indicates '2 instances selected' with a red box highlighting the 'START' button, along with 'RESET', 'STOP', and 'DELETE' buttons. The left sidebar contains navigation options: TOOLS (Dashboard, Model Garden, Pipelines), NOTEBOOKS (Colab Enterprise, Workbench), and VERTEX AI STUDIO (Overview, Multimodal, Language, Vision, Speech). The main content area shows a table of instances with columns: Instance name, Zone, Auto upgrade, Version, Machine Type, GPUs, Owner, Created, and Labels. Two instances are listed, both with checkboxes selected. The first instance is 'ai-sb-user-100-gpu' and the second is 'ai-sb-user-100-no-gpu'. Both are in the 'asia-southeast1-c' zone, version 'M118', and created on March 11, 2024.

Instance name	Zone	Auto upgrade	Version	Machine Type	GPUs	Owner	Created	Labels
<input checked="" type="checkbox"/> <a href="#">ai-sb-user-100-gpu</a>	asia-southeast1-c	—	M118	Graphics Optimized: 1 NVIDIA L4 GPU, 4 vCPUs, 16GB RAM	NVIDIA L4 x 1	55282892569-compute@developer.gserviceaccount.com	Mar 11, 2024, 7:49:27 AM	consumer-p...: ai-sandbox...
<input checked="" type="checkbox"/> <a href="#">ai-sb-user-100-no-gpu</a>	asia-southeast1-c	—	M118	Efficient Instance: 4 vCPUs, 16 GB RAM	None	55282892569-compute@developer.gserviceaccount.com	Mar 11, 2024, 7:56:19 AM	consumer-p...: ai-sandbox...

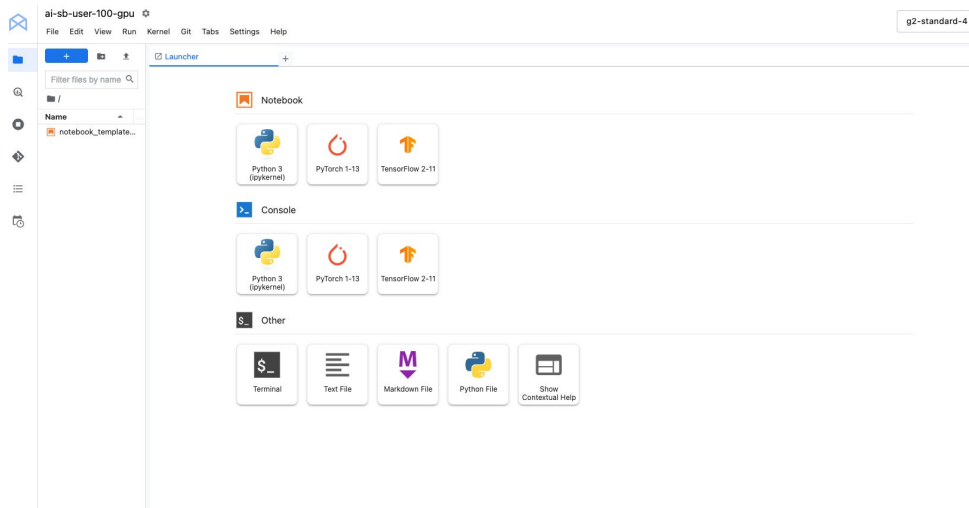
# Labs Setup: Accessing Workbench Notebook

**Step 7:** Wait 1-3 minutes for the instances to start. Verify that there is a green tick beside the instances.



Instance name	Zone	Auto upgrade	Version	Machine Type	GPUs	Owner
ai-sb-user-100-gpu	asia-southeast1-c	—	M118	Graphics Optimized: 1 NVIDIA L4 GPU, 4 vCPUs, 16GB RAM	NVIDIA L4 x 1	55282892569- compute@developer.gserviceaccount.com


**Step 8:** Click on **OPEN JUPYTERLAB** for the **gpu instance** <ai-gpu-100-XXX>. A new tab will open with access to the Jupyter Notebook.



# Labs Setup: Accessing Workbench Notebook

**Step 9:** Expand the **Git** menu and click on **Clone a Repository**

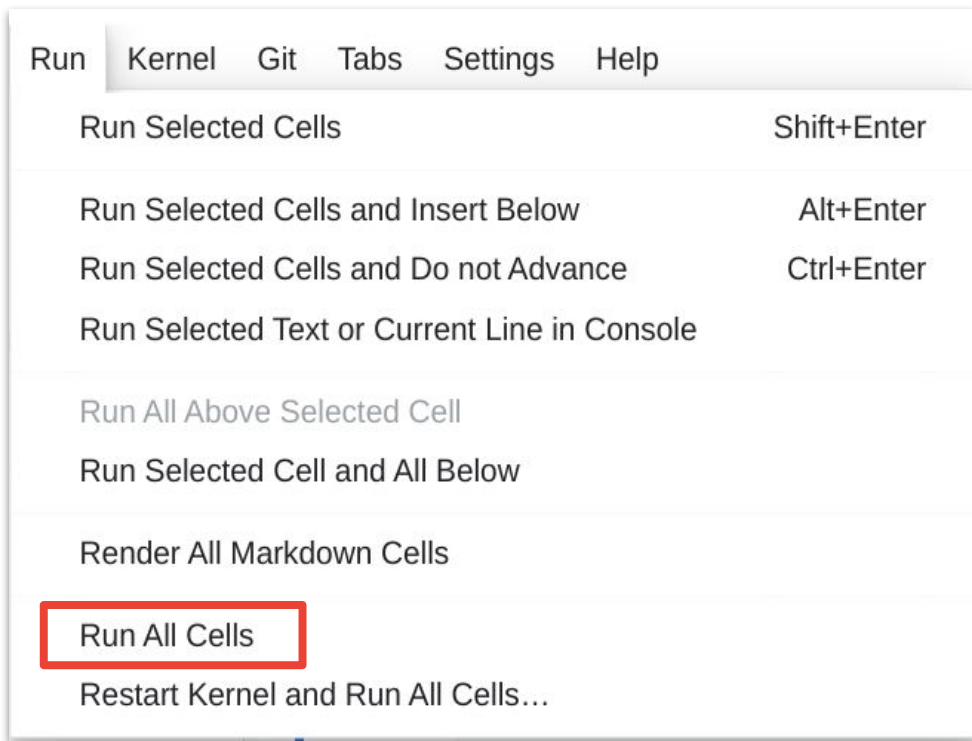
**Step 10:** Paste the URL provided for you into the text box and click **Clone**  
URL = `https://github.com/analyticsrepo01/pru-ai-labs.git`





# Labs Setup: Setting Up Notebooks

**Step 12:** Expand the **Run** menu and select **Run All Cells**








# Image model with custom training



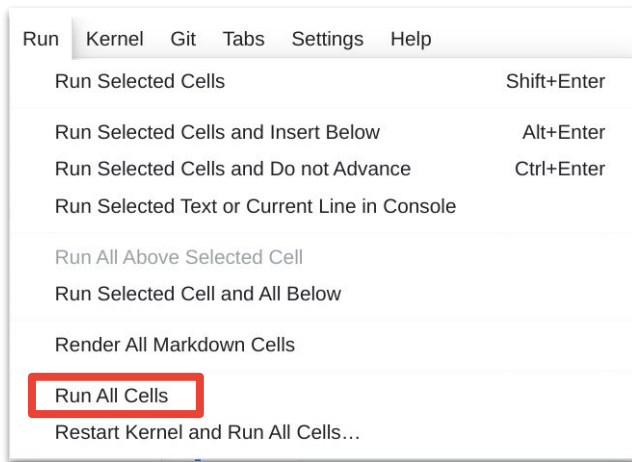
## Labs Setup: Setting Up Notebooks

**Step 11:** Double click [pru\\_sd\\_xl\\_finetuning\\_dreambooth\\_lora.ipynb](#) to open the next notebook

- my\_video.avi
-  part\_DIY\_RAG\_pdf.ipynb
-  part\_mixtral.ipynb
-  [pru\\_sd\\_xl\\_finetuning\\_dreambooth\\_lora.ipynb](#)
-  part5\_RAG\_VertexSearch.ipynb
-  part6\_Gemin\_on\_video.ipynb

# Labs Setup: Setting Up Notebooks

**Step 12:** Expand the **Run** menu and select **Run All Cells**.



# Open LLM with custom training



# Labs Setup: Tuning a Open source LLM

## Step 13: On Google Cloud Console, Navigate to Vertex AI → Model Garden

The screenshot displays the Google Cloud Console interface. At the top, the search bar contains 'Vertex AI', which is highlighted with a red rectangle. In the left-hand navigation menu, the 'Model Garden' option is also highlighted with a red rectangle. The main content area shows the 'Get started with Vertex AI' section, which includes a description of Vertex AI's capabilities and a button to 'ENABLE ALL RECOMMENDED APIS'. Below this, there are three cards: 'Vertex AI Studio', 'Model Garden', and 'Colab Enterprise', each with a 'Try now' or 'Go to Colab Enterprise' button. The 'Model Garden' card is the focus of the step.

Google Cloud AI-sandbox-Company118 Vertex AI Search

Vertex AI Dashboard

TOOLS

- Dashboard
- Model Garden**
- Pipelines

NOTEBOOKS

- Colab Enterprise
- Workbench

VERTEX AI STUDIO

- Overview
- Freeform
- Chat
- Vision
- Translation
- Speech
- Prompt gallery
- Prompt management
- Tuning
- Marketplace

**Get started with Vertex AI**

Vertex AI empowers machine learning developers, data scientists, and data engineers to take their projects from ideation to deployment, quickly and cost-effectively. [Learn more about Vertex AI](#)

[ENABLE ALL RECOMMENDED APIS](#)

[SHOW API LIST](#)

**Tutorials**

Learn how to use generative AI, Gemini, and more on Vertex AI.

[VIEW TUTORIALS](#)

**Vertex AI Studio**

Prompt, test and tune generative AI models like Gemini.

[Try now](#)

**Model Garden**

Browse, customize, and deploy machine learning models. Choose from Google or popular open-source models.

[Try now](#)

**Colab Enterprise**

A new notebook experience with enterprise-grade privacy and security. Start coding in a couple clicks.

[Go to Colab Enterprise](#)



## TOOLS



Dashboard



Model Garden



Pipelines

## NOTEBOOKS



Colab Enterprise



Workbench

## VERTEX AI STUDIO



Overview



Freeform



Chat



Vision



Translation



Speech



Prompt gallery



Prompt management



Tuning

## BUILD WITH GEN AI



Extensions



Marketplace

## Modalities

Language	67
Vision	88
Tabular	7
Document	8
Speech	2
Video	6
Multimodal	21
Audio	1

## Tasks

Generation	74
Classification	66
Detection	44
Extraction	28
Recognition	26
Translation	24
Embedding	7
Segmentation	12
Retrieval	2
Open vocabulary detection	2
Open vocabulary segmentation	2

Search models

Llama 3.1

Browse, customize, and deploy machine learning models with **Model Garden**. Choose from models created by Google and other providers.

Sort by: [Trending](#) [Newest](#) [Last Update](#)

## Foundation models

[SHOW ALL \(93\)](#)

Pre-trained multi-task models that can be further tuned or customized for specific tasks.



## Gemini 1.5 Pro

Created from the ground up to be multimodal (text, images, videos) and to scale across a wide range of tasks



## Gemini 1.5 Flash

The best performing Gemini model with features for a wide range of tasks



## Gemini 1.0 Pro



Designed to balance quality, performance, and cost for tasks such as content generation, editing, summarization, and classification



## Gemini 1.0 Pro Vision

Created to be multimodal (text, images, code) and to scale across a wide range of tasks

## Featured partners



# Step 14: Choose Llama 3.1 model, Click on “FINE-TUNE”

[←](#) Llama 3.1

## Llama 3.1

Explore and build with Llama 3.1 models on Vertex AI.

[DEPLOY](#)
[FINE-TUNE](#)
[CO OPEN NOTEBOOK](#)

[OVERVIEW](#)
[USE CASES](#)
[DOCUMENTATION](#)
[LICENSE](#)

★ Llama 3.1 is priced on a dollar-per-hour basis. You will need to deploy this model to Vertex AI. Alternatively, the Llama 3.1 API service is at no cost during public preview, and will be priced as per dollar-per-1M-tokens at GA. To learn more, see the [Llama 3.1 API service](#) model card.

### Overview

Note: Use of this model is governed by the Meta license. See the License tab.

The Meta Llama 3.1 collection of multilingual large language models (LLMs) is a collection of pretrained and instruction tuned generative models in 8B, 70B and 405B sizes (text in/text out). The Llama 3.1 instruction tuned text only models (8B, 70B, 405B) are optimized for multilingual dialogue use cases and outperform many of the available open source and closed chat models on common industry benchmarks.

**Model Developers** Meta

**Model Architecture** Llama 3.1 is an auto-regressive language model that uses an optimized transformer architecture. The tuned versions use supervised fine-tuning (SFT) and reinforcement learning with human feedback (RLHF) to align with human preferences for helpfulness and safety.

	Training data	Params	Input modalities	Output modalities	Content length	GQA	Tot cor
Llama 3.1 (text only)	A new mix of publicly available online data.	8B	Multilingual Text	Multilingual Text and Code	128k	Yes	15

### Try out Llama 3.1

Region \*  
us-central1 (Iowa)

Endpoint

To test this model, you will need to deploy it to a Vertex AI endpoint.

Prompt \*  
Text that describes what you want to generate

[ADVANCED OPTIONS](#)

SUBMIT

\* This demo is for internal testing purposes only. Output should not be saved or distributed. Please do not provide personally identifiable information or other data subject to regulatory requirements.

Model ID

Google



## Create a fine-tuned model

Fine tune this model using supervised tuning.

Tuned model name \*  
llama-3-1-70b-1731549015067

Base model  
Llama-3-1-70B

Output directory \* BROWSE ?

Input is required

Region \*  
us-central1 (Iowa)

The base model will be tuned with the following settings ([See pricing](#))

- Machine type: a2-ultragpu-8g
- Accelerator type: NVIDIA\_A100\_80GB
- Accelerator count: 8

### Tuning parameters

Number of epochs \*  
3

Learning rate \*  
0.0002

[VIEW TUNING CONFIG](#)

### Dataset

Tuning dataset \*  
☒ cloud-samples-data/vertex-ai/model-evaluation/peft\_train\_sam BROWSE ?

Evaluation dataset \*  
☒ cloud-samples-data/vertex-ai/model-evaluation/peft\_eval\_sam BROWSE ?

[VIEW EXAMPLE FORMAT](#)

START TUNING

CANCEL

Create or select the bucket

## Create a bucket

### Get Started

Pick a globally unique, permanent name. [Naming guidelines](#)

llama3trainingv1

Tip: Don't include any sensitive information

Optimize storage for data-intensive workloads

Labels (optional)

CONTINUE

### Choose where to store your data

Location: us (multiple regions in United States)

Location type: Multi-region

### Choose a storage class for your data

Default storage class: Standard

### Choose how to control access to objects

Public access prevention: On

Access control: Uniform

### Choose how to protect object data

Soft delete policy: Default

Object versioning: Disabled

Bucket retention policy: Disabled

Object retention: Disabled

Encryption type: Google-managed

CREATE

CANCEL

Next start tining once all filled up

Confidential

Vertex AI Training TRAIN NEW MODEL

TRAINING PIPELINES CUSTOM JOBS

Go to training tab - from left menu & check training started?

DATA	Extensions	<a href="#">llama3-1-lora-train-20241201-164565</a>	233993907094945792	Finished	Training pipeline	Custom	17 min 6 sec	Dec 2, 2024, 1:02:14 AM	Dec 2, 2024, 12:45:07 AM	Dec 2, 2024, 1:02:14 AM	—	⋮
	Code samples	<a href="#">llama31-lora-20241201-143742</a>	3612889896273838080	Finished	Training pipeline	Custom	28 min 8 sec	Dec 1, 2024, 11:05:52 PM	Dec 1, 2024, 10:37:44 PM	Dec 1, 2024, 11:05:52 PM	—	⋮
	Agents	<a href="#">llama3-1-lora-train-20241201-141259</a>	3170974182838108160	Failed	Training pipeline	Custom	52 min 45 sec	Dec 1, 2024, 11:05:45 PM	Dec 1, 2024, 10:13:00 PM	Dec 1, 2024, 11:05:45 PM	—	⋮
	Feature Store	<a href="#">llama31-lora-20241201-143729-1733063850</a>	2477982790176473088	Finished	Training pipeline	Custom	28 min 8 sec	Dec 1, 2024, 11:05:39 PM	Dec 1, 2024, 10:37:30 PM	Dec 1, 2024, 11:05:39 PM	—	⋮
MODEL DEVELOPMENT	Datasets	<a href="#">llama3-1-lora-train-20241201-144150</a>	3286941873242898432	Finished	Training pipeline	Custom	23 min 7 sec	Dec 1, 2024, 11:04:59 PM	Dec 1, 2024, 10:41:52 PM	Dec 1, 2024, 11:04:59 PM	—	⋮
	Training	<a href="#">llama3-1-lora-train-20241201-134300</a>	3813018604715114496	Failed	Training pipeline	Custom	52 min 14 sec	Dec 1, 2024, 10:35:16 PM	Dec 1, 2024, 9:43:01 PM	Dec 1, 2024, 10:35:16 PM	—	⋮
	Experiments	<a href="#">llama31-lora-20241201-143018</a>	478384555623972864	Failed	Training pipeline	Custom	1 min 32 sec	Dec 1, 2024, 10:31:52 PM	Dec 1, 2024, 10:30:20 PM	Dec 1, 2024, 10:31:52 PM	—	⋮
	Metadata	<a href="#">llama3-1-lora-train-20241201-141611</a>	4672924658566168576	Finished	Training pipeline	Custom	15 min 34 sec	Dec 1, 2024, 10:31:47 PM	Dec 1, 2024, 10:16:12 PM	Dec 1, 2024, 10:31:47 PM	—	⋮
DEPLOY AND USE	Model Registry	<a href="#">llama3-1-lora-train-20241201-141539</a>	5396878298665975808	Finished	Training pipeline	Custom	15 min 35 sec	Dec 1, 2024, 10:31:16 PM	Dec 1, 2024, 10:15:40 PM	Dec 1, 2024, 10:31:16 PM	—	⋮
	Online prediction	<a href="#">llama3-1-lora-train-20241201-141534</a>	4221438795922276352	Finished	Training pipeline	Custom	15 min 35 sec	Dec 1, 2024, 10:31:11 PM	Dec 1, 2024, 10:15:35 PM	Dec 1, 2024, 10:31:11 PM	—	⋮
	Batch predictions	<a href="#">llama3-1-lora-train-20241201-142214</a>	1870559790434877440	Failed	Training pipeline	Custom	5 min 2 sec	Dec 1, 2024, 10:27:19 PM	Dec 1, 2024, 10:22:16 PM	Dec 1, 2024, 10:27:19 PM	—	⋮
	Monitoring	<a href="#">llama3-1-lora-train-20241201-141756</a>	9148376788265598976	Failed	Training pipeline	Custom	5 min 33 sec	Dec 1, 2024, 10:23:32 PM	Dec 1, 2024, 10:17:58 PM	Dec 1, 2024, 10:23:32 PM	—	⋮
MANAGE	Vector Search	<a href="#">llama3-1-lora-train-20241201-131514</a>	6548955378342690816	Failed	Training pipeline	Custom	53 min 14 sec	Dec 1, 2024, 10:08:31 PM	Dec 1, 2024, 9:15:16 PM	Dec 1, 2024, 10:08:31 PM	—	⋮
	Ray on Vertex AI	<a href="#">llama3-1-lora-train-20241201-134445</a>	6964412443967619072	Finished	Training pipeline	Custom	16 min 35 sec	Dec 1, 2024, 10:01:23 PM	Dec 1, 2024, 9:44:47 PM	Dec 1, 2024, 10:01:23 PM	—	⋮
	Migrate to Vertex AI	<a href="#">llama3-1-lora-train-20241201-134438</a>	3384050740208074752	Finished	Training pipeline	Custom	16 min 5 sec	Dec 1, 2024, 10:00:46 PM	Dec 1, 2024, 9:44:40 PM	Dec 1, 2024, 10:00:46 PM	—	⋮
	Marketplace	<a href="#">llama3-1-lora-train-20241201-132545</a>	2620622010806600408	Finished	Training pipeline	Custom	13 min 4 sec	Dec 1, 2024, 9:38:59 PM	Dec 1, 2024, 9:25:47 PM	Dec 1, 2024, 9:38:59 PM	—	⋮

# Example on custom data: Llama3.1 on Insurance data

Open and run all the notebook [pru\\_llama3\\_1\\_fineting.ipynb](#)

The screenshot shows a Google Cloud Vertex AI notebook interface. On the left, a file explorer shows the directory structure under 'GenAI8 / pru-ai-labs /'. The file 'pru\_llama3\_1\_fineting.ipynb' is selected. The main area displays the notebook content, which includes a title 'Vertex AI Model Garden - Llama 3.1 Fineting Insurance Data', an overview, objective, costs, and before you begin sections.

**Vertex AI Model Garden - Llama 3.1 Fineting Insurance Data**

Run in Colab Enterprise View on GitHub

**Overview**

This notebook demonstrates finetuning and deploying Llama 3.1 models with Vertex AI. After finetuning, we can deploy models on Vertex with GPU.

**Objective**

- Finetune Llama 3.1 models with Vertex AI Custom Training Jobs.
- Deploy finetuned Llama 3.1 models on Vertex AI Prediction.
- Send prediction requests to your finetuned Llama 3.1 models.

**Costs**

This tutorial uses billable components of Google Cloud:

- Vertex AI
- Cloud Storage

**Before you begin**

**Install dependencies**

```
[21]: 1 print("Installing google-cloud-aiplatform")
```

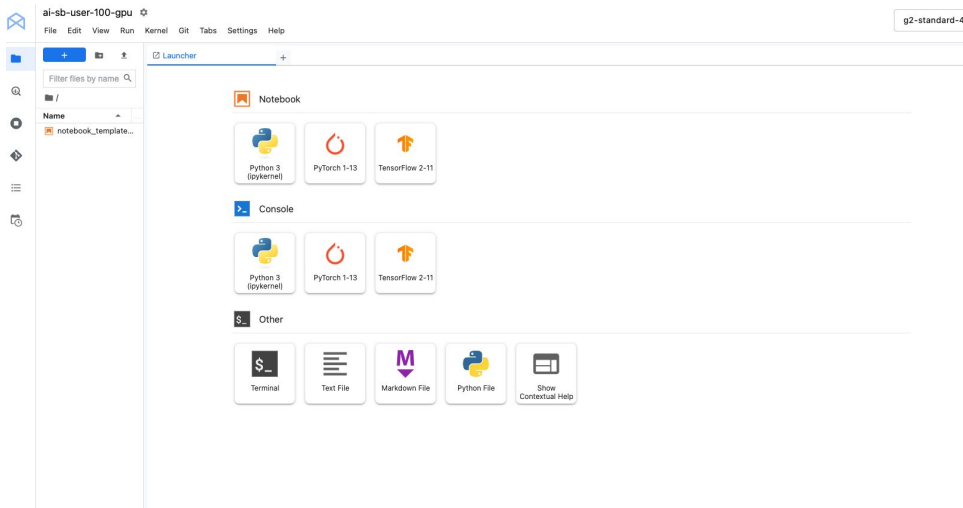


# RAG for Prudential Use case



# Labs Setup: Accessing Workbench Notebook

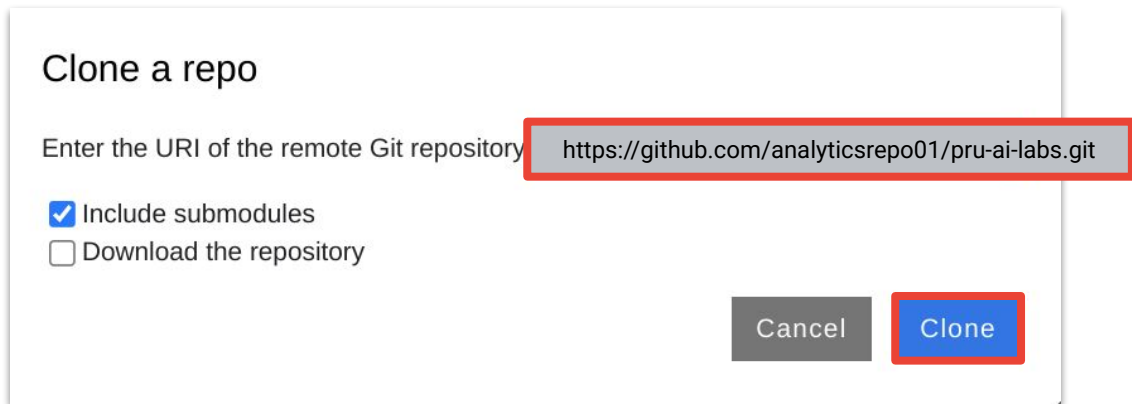
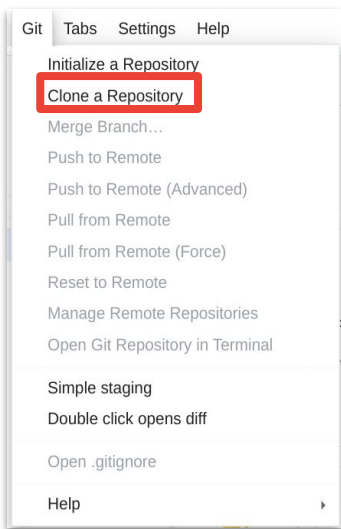
**Step 15:** Click on **OPEN JUPYTERLAB** for the **no-gpu** instance <ai-no-gpu-100-XXX>. A new tab will open with access to the Jupyter Notebook.



# Labs Setup: Accessing Workbench Notebook

**Step 16:** Expand the **Git** menu and click on **Clone a Repository**

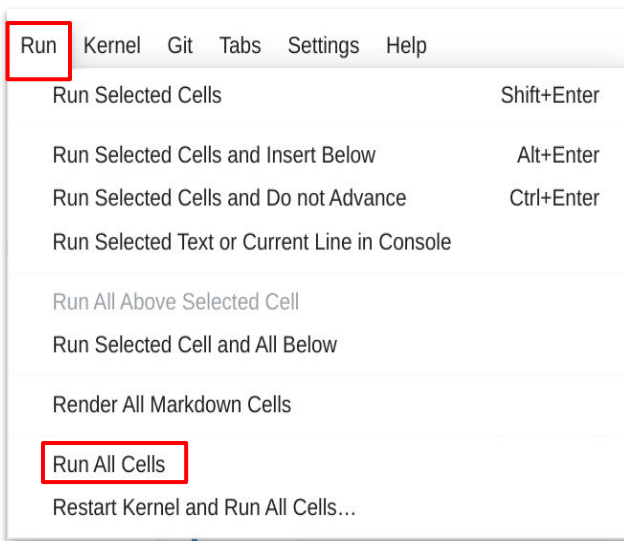
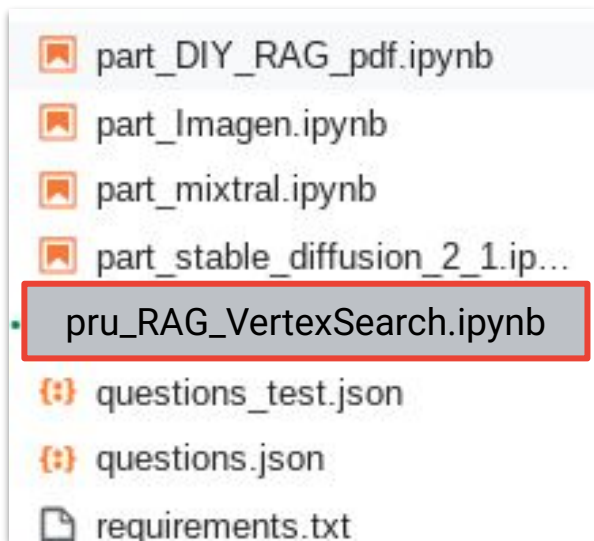
**Step 17:** Paste the URL provided for you into the text box and click **Clone**  
URL = `https://github.com/analyticsrepo01/pru-ai-labs.git`



# Labs Setup: Vertex AI Search

**Step 18:** Double click **pru\_RAG\_VertexSearch.ipynb** to open the Jupyter Notebook

**Step 19:** In the top menu bar, expand the **Run** menu and select **Run All Cells**






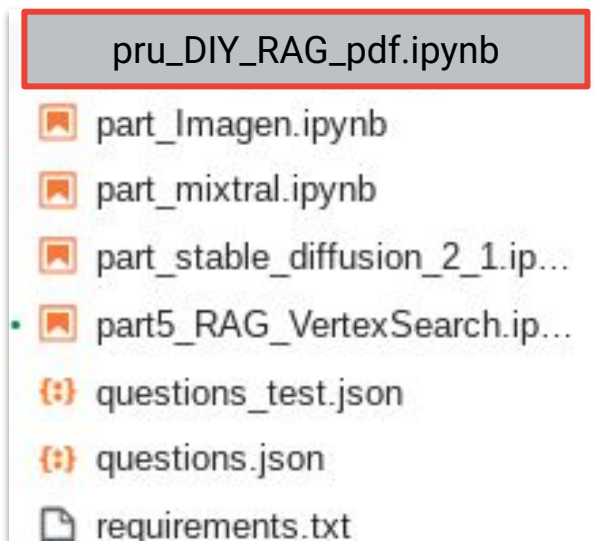
# DIY RAG



# Labs Setup: Vector AI Search

**Step 20:** Double click `pru_DIY_RAG_pdf.ipynb` to open the Jupyter Notebook

**Step 21:** In the top menu bar, expand the **Run** menu and select **Run All Cells**



# Kernel Restarts after installation

Step1 : Click on this cell after restart

The screenshot shows a Google Colab notebook interface. The top bar displays the URL `6f11ba225292928c-dot-us-central1.notebooks.googleusercontent.com/lab/tree/pru-ai-labs/pru_DIY_RAG_pdf.ipynb`. The notebook is titled `ai-no-gpu-100-e3d5fc22` and is running on `n1-standard-4`. The left sidebar shows a file explorer with a tree view of the notebook's contents, including folders like `backup_folder`, `books`, `campaign`, `dog`, `images`, `results`, `testing`, `vertex-ai-samples`, `video`, and `web-app`. The main area shows a code cell with the following content:

```
[1]: ## Run this cell and all below

Declare Variables

[2]: #no spaces or special characters allowed), ensure that it is unique
import socket
import re

UNIQUE_PREFIX = socket.gethostname()

UNIQUE_PREFIX = re.sub('[^A-Za-z0-9]+', '', UNIQUE_PREFIX)

PREFIX_NUM_ONLY=int(str(re.search(r'\d+', UNIQUE_PREFIX).group()))

REGION_ALLOCATE=PREFIX_NUM_ONLY%3

[3]: PROJECT_ID = !(gcloud config get-value core/project)
PROJECT_ID = PROJECT_ID[0]

SVC_ACC = !(gcloud config get-value core/account)
SVC_ACC = SVC_ACC[0]

PROJECT_NUMBER=str(re.search(r'\d+', SVC_ACC).group())

LOCATION="asia-southeast1"

FOLDER_NAME=""
```

The notebook also shows a list of recent files in the sidebar, including `pru_DIY_RAG_pdf.ipynb`, `pru_llama3_1_finetuning.ipynb`, `pru_qa.csv`, `pru_RAG_VertexSearch.ipynb`, `pru_sd_xl_finetuning_dreambooth_lora...`, `pru_stable_diffusion_2_1.ipynb`, and `questions_test.json`.

Step2 : Run  
Selected Cell  
and All Below

# Multi Agents



ai-gpu-100-f20436e7

g2-standard-4

File Edit View Run Kernel Git Tabs Settings Help



Filter files by name	
/ pru-ai-labs /	
Name	Last Modified
images	a day ago
results	a day ago
testing	a day ago
vertex-ai-samples	a day ago
video	a day ago
web-app	a day ago
\$BUCKET_URI	a day ago
Agent_Email_conversation002.ip...	2 minutes ago
Agent_Example.ipynb	a day ago
Imagen_on_questions.ipynb	a day ago
intro_multimodal_use_cases_lat...	a day ago
mdpdf.log	6 minutes ago
Multi-agent_002.ipynb	seconds ago
part_DIY_RAG_pdf.ipynb	a day ago
part_Imagen.ipynb	a day ago
part_llama3.ipynb	a day ago
part_mixtral.ipynb	a day ago
part6_Gemini_on_video.ipynb	a day ago
pru_llama3_1_finetuning.ipynb	4 hours ago
pru_qa.csv	a day ago
pru_RAG_VertexSearch.ipynb	a day ago
pru_sd_xl_finetuning_dreamboo...	a day ago
pru_stable_diffusion_2_1.ipynb	a day ago
questions_test.json	a day ago
questions.json	a day ago

Terminal 1

pru\_llama3\_1\_finetuning.ipynb Agent\_Email\_conversationC Multi-agent\_002.ipynb

Code git Execute

file\_read\_tool 1/2

PyTorch 1-13 (Local)

```
[1]: # !conda create -n crewai python=3.11
# !conda activate crewai -
!pip install -q --upgrade google-cloud-aiplatform
!pip install -q -U 'crewai[tools]' mdpdf

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
gcsfs 2024.10.0 requires fsspec==2024.10.0, but you have fsspec 2024.9.0 which is incompatible.
kfp 2.5.0 requires kubernetes<27,>=8.0.0, but you have kubernetes 31.0.0 which is incompatible.
kfp 2.5.0 requires requests-toolbelt<1,>=0.8.0, but you have requests-toolbelt 1.0.0 which is incompatible.
kfp 2.5.0 requires urllib3<2.0.0, but you have urllib3 2.2.3 which is incompatible.
```

```
[2]: import re

PROJECT_ID = !(gcloud config get-value core/project)
PROJECT_ID = PROJECT_ID[0]

SVC_ACC = !(gcloud config get-value core/account)
SVC_ACC = SVC_ACC[0]

PROJECT_NUMBER=str(re.search(r'\d+', SVC_ACC).group())

LOCATION="asia-southeast1"

FOLDER_NAME="."
```

```
[3]: from crewai import Agent, Task, Crew, Process
from crewai_tools import tool
from langchain_openai import ChatOpenAI
from crewai_tools.tools import FileReadTool
import os, requests, re, mdpdf, subprocess
from openai import OpenAI
```

```
[4]: !pip install --upgrade --quiet langchain-core langchain-google-vertexai
!pip install mdpdf

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
```

## Optional Lab :

Go to this github URL and clone

[https://github.com/analyticsrepo01/ai\\_agents\\_v2](https://github.com/analyticsrepo01/ai_agents_v2)

`git clone <URL>`



## Labs Setup: Setting Up Notebooks

### Note:

**Avoid letting your laptops/computers enter sleep mode to prevent problems arising when running the notebooks**

# Thank you





# Multi-Agent Lab: Generating AI News Insights

## Objectives:

- Retrieve news about AI using RAG
- Generate summarized reports on AI trends

